

Strategy for developing efficient programs:

1. Design the program well
2. Implement the program well**
3. Test the program well
4. Only after you're sure it's working, *measure* performance
5. If (and only if) performance is inadequate, *find* the "hot spots"
6. *Tune* the code to fix these
7. Repeat measure-analyse-tune cycle until performance ok

(** see "Programming Pearls", "Practice of Programming", etc. etc.)

Rapid development of a prototype may be the best way to discover/assess performance issues.

Hence Fred Brooks maxim - "Plan To Throw One Away".

simple C program which allocate+initialializes an array

/home/cs2041/public_html/code/performance/cachegrind_example.c

```
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>

void test0(int x, int y, int a[x][y]) {
    int i, j;
    fprintf(stderr, "writing to array i-j order\n");
    for (i = 0; i < x; i++)
        for (j = 0; j < y; j++)
            a[i][j] = i+j;
}

void test1(int x, int y, int a[x][y]) {
    fprintf(stderr, "writing to array j-i order\n");
    int i, j;
    for (j = 0; j < y; j++)
        for (i = 0; i < x; i++)
            a[i][j] = i+j;
}

int main(int argc, char*argv[]) {
    int x = atoi(argv[2]);
    int y = atoi(argv[3]);
    fprintf(stderr, "allocating a %dx%d array = %lld bytes\n", x, y, ((long long)x)*y*sizeof (int));
    void *m = malloc(x*y*sizeof (int));
    assert(m);
    switch (atoi(argv[1])) {
        case 0: test0(x, y, m); break;
        case 1: test1(x, y, m); break;
    }
    return 0;
}
```

Performance Improvement Example - Caching

Although the loops are almost identical, the first loop runs 20x faster on a large array!

```
$ time ./cachegrind_example 0 32000 32000
allocating a 32000x32000 array = 4096000000 bytes
writing to array i-j order
real    0m0.893s
user    0m0.364s
sys     0m0.524s
$ time ./cachegrind_example 1 32000 32000
allocating a 32000x32000 array = 4096000000 bytes
writing to array j-i order
real    0m15.189s
user    0m14.633s
sys     0m0.528s
```

Performance Improvement Example - Caching

The tool valgrind used to detect accesses to uninitialized variables at runtime also can give memory caching information.

The memory subsystem is beyond the scope of this course but you can see valgrind explain the performance difference between these loops.

For the first loop D1 miss rate = 24.8%

For the second loop D1 miss rate = 99.9%

Due to the C array memory layout the first loop produces much better caching performance.

Tuning caching performance is important for some application and valgrind makes this much easier.

Performance Improvement Example - Caching

```
$ valgrind '--tool=cachegrind' ./cachegrind_example 0 10000 10000
allocating a 10000x10000 array = 400000000 bytes
writing to array i-j order
==7025==
==7025== I   refs:      225,642,966
==7025== I1  misses:      882
==7025== LLi misses:      875
==7025== I1  miss rate:    0.00%
==7025== LLi miss rate:    0.00%
==7025==
==7025== D   refs:      25,156,289 (93,484 rd + 25,062,805 wr)
==7025== D1  misses:      6,262,957 ( 2,406 rd +  6,260,551 wr)
==7025== LLd misses:      6,252,482 ( 1,982 rd +  6,250,500 wr)
==7025== D1  miss rate:    24.8% (  2.5% +    24.9% )
==7025== LLd miss rate:    24.8% (  2.1% +    24.9% )
==7025==
==7025== LL refs:      6,263,839 ( 3,288 rd +  6,260,551 wr)
==7025== LL misses:      6,253,357 ( 2,857 rd +  6,250,500 wr)
==7025== LL miss rate:     2.4% (  0.0% +    24.9% )
```

Performance Improvement Example - Caching

```
$ valgrind '--tool=cachegrind' ./cachegrind_example 1 10000 10000
allocating a 10000x10000 array = 400000000 bytes
writing to array j-i order
==7006==
==7006== I   refs:      600,262,960
==7006== I1  misses:      876
==7006== LLi misses:      869
==7006== I1  miss rate:    0.00%
==7006== LLi miss rate:    0.00%
==7006==
==7006== D   refs:      100,056,288 (43,483 rd + 100,012,805 wr)
==7006== D1  misses:      100,002,957 ( 2,405 rd + 100,000,552 wr)
==7006== LLd misses:      6,262,481 ( 1,982 rd +  6,260,499 wr)
==7006== D1  miss rate:    99.9% (  5.5% +    99.9% )
==7006== LLd miss rate:     6.2% (  4.5% +     6.2% )
==7006==
==7006== LL refs:      100,003,833 ( 3,281 rd + 100,000,552 wr)
==7006== LL misses:      6,263,350 ( 2,851 rd +  6,260,499 wr)
==7006== LL miss rate:     0.8% (  0.0% +     6.2% )
```

Where is execution time being spent?

Typically programs spend most of their execution time in a small part of their code.

This is often quoted as the 90/10 rule (or 80/20 rule or ...):

“90% of the execution time is spent in 10% of the code”

This means that

- most of the code has little impact on overall performance
- small parts of the code account for most execution time

We should clearly concentrate efforts at improving execution speed in the 10% of code which accounts for most of the execution time.

gcc -p/gprof

Given the -p flag gcc instruments a C program to collect profile information

When the program executes this data is left in the file gmon.out.

The program gprof analyzes this data and produces:

- number of times each function was called
- % of total execution time spent in the function
- average execution time per call to that function
- execution time for this function and its children

Arranged in order from most expensive function down.

It also gives a *call graph*, a list for each function:

- which functions called this function
- which functions were called by this function

Table 1 Summary of the data sets used in the study. The data sets are categorized into three groups: *General*, *Specific*, and *Specialized*. The *General* group includes data sets for general text classification, the *Specific* group includes data sets for specific tasks, and the *Specialized* group includes data sets for specialized tasks. The data sets are described in terms of their source, size, and the tasks they are used for.

Category	Dataset	Source	Size	Task
General	1.1	General	100,000	Text Classification
	1.2	General	50,000	Text Classification
	1.3	General	20,000	Text Classification
	1.4	General	10,000	Text Classification
Specific	2.1	Specific	5,000	Text Classification
	2.2	Specific	2,000	Text Classification
	2.3	Specific	1,000	Text Classification
	2.4	Specific	500	Text Classification
Specialized	3.1	Specialized	1,000	Text Classification
	3.2	Specialized	500	Text Classification
	3.3	Specialized	250	Text Classification
	3.4	Specialized	125	Text Classification

```
Program is slow on large inputs e.g.
```

```
$ gcc -O3 word_frequency0.c -o word_frequency0
$ time word_frequency0 <WarAndPeace.txt >/dev/null
real    0m52.726s
user    0m52.643s
sys     0m0.020s
```

```
We can instrument the program to collect profiling information and
examine it with gcc

$ gcc -p -g word_frequency0.c -o word_frequency0_profile
$ head -10000 WarAndPeace.txt|word_frequency0_profile >/dev/null
$ gprof word_frequency0_profile

....
%   cumulative    self           self         total
time  seconds    seconds   calls   ms/call  ms/call  name
88.90      0.79      0.79    88335     0.01     0.01   get
 7.88      0.86      0.07     7531     0.01     0.01   put
 2.25      0.88      0.02    80805     0.00     0.00   get_word
 1.13      0.89      0.01         1    10.02    823.90   read_words
 0.00      0.89      0.00         2     0.00     0.00   size
 0.00      0.89      0.00         1     0.00     0.00   create_map
 0.00      0.89      0.00         1     0.00     0.00   keys
 0.00      0.89      0.00         1     0.00     0.00   sort_words
....
```

put are relevant to performance improvement.

Examine *get* and we find it traverses a linked list.
So replace it with a binary tree and the program runs 200x faster on War and Peace.

```
$ gcc -O3 word_frequency1.c -o word_frequency1
$ time word_frequency1 <WarAndPeace.txt >/dev/null
real    0m0.277s
user    0m0.268s
sys     0m0.008s
```

Was C the best choice for our count words program.

Shell - Word Count

For comparison Shell:

/home/cs2041/public_html/code/performance/word_frequency.sh

```
#!/bin/sh
tr -c a-zA-Z ' '|
tr ' ' '\n'|
tr A-Z a-z|
egrep -v '^$'|
sort|
uniq -c
```

Perl - Word Count

/home/cs2041/public_html/code/performance/word_frequency.pl

```
#!/usr/bin/perl -w

while ($line = <>) {
    $line =~ tr/A-Z/a-z/;
    foreach $word ($line =~ /[a-z]+/g) {
        $count{$word}++;
    }
}

@words = keys %count;

@sorted_words = sort {$count{$a} <=> $count{$b}} @words;

foreach $word (@sorted_words) {
    printf "%8d %s\n", $count{$word}, $word;
}
```

Python - Word Count

/home/cs2041/public_html/code/performance/word_frequency.py

```
#!/usr/bin/python3
import fileinput, re, collections

count = collections.defaultdict(int)
for line in fileinput.input():
    for word in re.findall(r'\w+', line.lower()):
        count[word] += 1

words = count.keys()

sorted_words = sorted(words, key=lambda w: count[w])

for word in sorted_words:
    print("%8d %s" % (count[word], word))
```

Performance Improvement Example - Word Count

Shell, Perl and Python are slower - but a lot less code.
So faster to write, less bugs to find, easier to maintain/modify

```
$ time word_frequency1 <WarAndPeace.txt >/dev/null
real    0m0.277s
user    0m0.268s
sys     0m0.008s
$ time word_frequency.sh <WarAndPeace.txt >/dev/null
real    0m0.564s
user    0m0.584s
sys     0m0.036s
$ time word_frequency.pl <WarAndPeace.txt >/dev/null
real    0m0.643s
user    0m0.632s
sys     0m0.012s
$ time word_frequency.py <WarAndPeace.txt >/dev/null
real    0m1.046s
user    0m0.836s
sys     0m0.024s
$ wc word_frequency1.c word_frequency.sh word_frequency.pl
286  759 5912 word_frequency1.c
  8   19   82 word_frequency.sh
11   38  325 word_frequency.py
14   43  301 word_frequency.pl
```

Performance Improvement Example - cp

```
Here is a cp implementation in C using low-level calls to read/write
/home/cs2041/public_html/codes/performance/cp.c

// Written by andrew@cse.unsw.edu.au
// as a COMP2041 lecture example

// copy input to output using read/write system calls
// for each byte - very inefficient and Unix/Linux specific

#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <stdlib.h>
#include <unistd.h>

void
copy_file_to_file(int in_fd, int out_fd) {
    while (1) {
        char c[1];
        int bytes_read = read(in_fd, c, 1);
        if (bytes_read < 0) {
            perror("cp: ");
            exit(1);
        }
        if (bytes_read == 0)
            return;
        int bytes_written = write(out_fd, c, bytes_read);
        if (bytes_written <= 0) {
            perror("cp: ");
            exit(1);
        }
    }
}

int
main(int argc, char *argv[]) {
    if (argc != 3) {
        fprintf(stderr, "cp <src-file> <destination-file>\n");
        return 1;
    }

    int in_fd = open(argv[1], O_RDONLY);
    if (in_fd < 0) {
        fprintf(stderr, "cp: %s: ", argv[1]);
        perror("");
        return 1;
    }

    int out_fd = open(argv[2], O_WRONLY|O_CREAT|O_TRUNC, S_IRWXU);
    if (out_fd <= 0) {
        fprintf(stderr, "cp: %s: ", argv[2]);
        perror("");
        return 1;
    }

    copy_file_to_file(in_fd, out_fd);
    return 0;
}
```

Performance Improvement Example - cp

Its suprisingly slow compared to /bin/cp

```
$ time /bin/cp input_file /dev/null
real    0m0.006s
user    0m0.000s
sys     0m0.004s
$ gcc cp0.c -o cp0
$ time ./cp0 input_file /dev/null
real    0m6.683s
user    0m0.932s
sys     0m5.740s
$ gcc -O3 cp0.c -o cp0
$ time ./cp0 input_file /dev/null
real    0m6.688s
user    0m0.900s
sys     0m5.776s
```

Performance Improvement Example - cp

Notice that most execution time is spent executing system calls and as a consequence gcc -O3 is no help.

cp0.c is making 2 system calls for every byte copied - a huge overhead.

If it is modified to buffer its I/O into 8192 byte block it runs 1000x faster.

It now makes 2 system calls for every 8192 bytes copied

Performance Improvement Example - cp

```
/home/cs2041/public_html/codes/performance/cp.c

// Written by andrew@cse.unsw.edu.au
// as a COMP2041 lecture example

// copy input to output using read/write system calls
// for every 4096 bytes - efficient but Unix/Linux specific

#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <stdlib.h>
#include <unistd.h>

void
copy_file_to_file(int in_fd, int out_fd) {
    while (1) {
        char c[8192];
        int bytes_read = read(in_fd, c, sizeof c);
        if (bytes_read < 0) {
            perror("cp: ");
            exit(1);
        }
        if (bytes_read <= 0)
            return;
        int bytes_written = write(out_fd, c, bytes_read);
        if (bytes_written <= 0) {
            perror("cp: ");
            exit(1);
        }
    }
}

int
main(int argc, char *argv[]) {
    if (argc != 3) {
        fprintf(stderr, "cp <src-file> <destination-file>\n");
        return 1;
    }

    int in_fd = open(argv[1], O_RDONLY);
    if (in_fd < 0) {
        fprintf(stderr, "cp: %s: ", argv[1]);
        perror("");
        return 1;
    }

    int out_fd = open(argv[2], O_WRONLY|O_CREAT|O_TRUNC, S_IRWXU);
    if (out_fd <= 0) {
        fprintf(stderr, "cp: %s: ", argv[2]);
        perror("");
        return 1;
    }

    copy_file_to_file(in_fd, out_fd);
    return 0;
}
```

Performance Improvement Example - cp

```
$ time ./cp1 input_file /dev/null
real    0m0.005s
user    0m0.000s
sys     0m0.008s
```

If we use the portable stdio library instead of a low-level copy, even a byte-byte by loop runs suprisingly fast, because stdio buffers the I/O behind the scenes.

Performance Improvement Example - cp

```
/home/cv2041/public_html/code/performance/cp2.c
// Written by andrew@cse.unsw.edu.au
// as a COMP2041 lecture example

// copy input to output using stdio functions
// stdio buffers reads & writes for you - efficient and portable

#include <stdio.h>
#include <stdlib.h>

void
copy_file_to_file(FILE *in, FILE *out) {
    while (!) {
        int ch = fgetc(in);
        if (ch == EOF)
            break;
        if (fputc(ch, out) == EOF) {
            fprintf(stderr, "cp: ");
            perror("");
            exit(1);
        }
    }
}

int
main(int argc, char *argv[]) {
    FILE *in, *out;

    if (argc != 3) {
        fprintf(stderr, "cp <src-file> <destination-file>\n");
        return 1;
    }

    in = fopen(argv[1], "r");
    if (in == NULL) {
        fprintf(stderr, "cp: %s: ", argv[1]);
        perror("");
        return 1;
    }

    out = fopen(argv[2], "w");
    if (out == NULL) {
        fprintf(stderr, "cp: %s: ", argv[2]);
        perror("");
        return 1;
    }

    copy_file_to_file(in, out);
    return 0;
}
```

Performance Improvement Example - cp

```
$ gcc -O3 cp2.c -o cp2
$ time ./cp2 input_file /dev/null
real    0m0.456s
user    0m0.448s
sys     0m0.008s
```

And with a little more complex code we get reasonable speed with portability:

Performance Improvement Example - cp

```
/home/cv2041/public_html/code/performance/cp3.c
// Written by andrew@cse.unsw.edu.au
// as a COMP2041 lecture example

#include <stdio.h>
#include <stdlib.h>

// copy input to output using stdio functions
// stdio buffers reads & writes for you - efficient and portable

void
copy_file_to_file(FILE *in, FILE *out) {
    char input[4096];

    while (!) {
        if (fgets(input, sizeof input, in) == NULL) {
            break;
        }
        if (fprintf(out, "%s", input) == EOF) {
            fprintf(stderr, "cp: ");
            perror("");
            exit(1);
        }
    }
}

int
main(int argc, char *argv[]) {
    FILE *in, *out;

    if (argc != 3) {
        fprintf(stderr, "cp <src-file> <destination-file>\n");
        return 1;
    }

    in = fopen(argv[1], "r");
    if (in == NULL) {
        fprintf(stderr, "cp: %s: ", argv[1]);
        perror("");
        return 1;
    }

    out = fopen(argv[2], "w");
    if (out == NULL) {
        fprintf(stderr, "cp: %s: ", argv[2]);
        perror("");
        return 1;
    }

    copy_file_to_file(in, out);
    return 0;
}
```

Performance Improvement Example - cp

```
$ gcc -O3 cp3.c -o cp3
$ time ./cp7 input_file /dev/null
real    0m0.095s
user    0m0.084s
sys     0m0.012s
```

Performance Improvement Example - cp

For comparison Perl code which does a copy via an array of lines:

/home/cs2041/public_html/code/performance/cp4.pl

```
#!/usr/bin/perl -w
# Simple cp implementation reading entire file into array
# Written by andrewt@cse.unsw.edu.au for COMP2041

die "Usage: cp <infile> <outfile>\n" if @ARGV != 2;
$infile = shift @ARGV;
$outfile = shift @ARGV;
open IN, '<', $infile or die "Cannot open $infile: $!\n";
open OUT, '>', $outfile or die "Cannot open $outfile: $!\n";
print OUT <IN>;

real    0m0.248s
user    0m0.168s
sys     0m0.032s
```

Performance Improvement Example - cp

And Perl code which unsets Perl's line terminator variable so a single read returns the whole file:

/home/cs2041/public_html/code/performance/cp5.pl

```
#!/usr/bin/perl -w
# Simple cp implementation reading entire file into array
# Written by andrewt@cse.unsw.edu.au for COMP2041

die "Usage: cp <infile> <outfile>\n" if @ARGV != 2;
$infile = shift @ARGV;
$outfile = shift @ARGV;
open IN, '<', $infile or die "Cannot open $infile: $!\n";
open OUT, '>', $outfile or die "Cannot open $outfile: $!\n";
undef $/;
print OUT <IN>;

real    0m0.029s
user    0m0.008s
sys     0m0.020s
```

Performance Improvement Example - Fibonacci

Here is a simple Perl program to calculate the n-th Fibonacci number:

/home/cs2041/public_html/code/performance/fib0.pl

```
#!/usr/bin/perl -w
sub fib($);
printf "fib(%d) = %d\n", $_, fib($_) foreach @ARGV;
sub fib($) {
    my ($n) = @_;
    return 1 if $n < 3;
    return fib($n-1) + fib($n-2);
}
```

Performance Improvement Example - Fibonacci

It becomes slow near $n=35$.

```
$ time fib0.pl 35
fib(35) = 9227465
real    0m10.776s
user    0m10.729s
sys     0m0.016s
```

we can rewrite in C.

Performance Improvement Example - Fibonacci

/home/cs2041/public_html/code/performance/fib0.c

```
#include <stdlib.h>
#include <stdio.h>

int fib(int n) {
    if (n < 3) return 1;
    return fib(n-1) + fib(n-2);
}

int main(int argc, char *argv[]) {
    int i;
    for (i = 1; i < argc; i++) {
        int n = atoi(argv[i]);
        printf("fib(%d) = %d\n", n, fib(n));
    }
    return 0;
}
```

Performance Improvement Example - Fibonacci

Faster but the program's complexity doesn't change and it becomes become slow near $n=45$.

```
$ gcc -O3 -o fib0 fib0.c
$ time fib0 35
fib(45) = 1134903170
real    0m4.994s
user    0m4.976s
sys     0m0.004s
```

It is very easy to cache already computed results in a Perl hash.

Performance Improvement Example - Fibonacci

/home/cs2041/public_html/code/performance/fib1.pl

```
#!/usr/bin/perl -w
sub fib($);
printf "fib(%d) = %d\n", $_, fib($_) foreach @ARGV;
sub fib($) {
    my ($n) = @_;
    return 1 if $n < 3;
    return $fib{$n} || ($fib{$n} = fib($n-1) + fib($n-2));
}
```


Performance Improvement Example - Fibonacci

This changes the program's complexity from exponential to linear.

```
$ time fib1.pl 45
fib(45) = 1134903170
real    0m0.004s
user    0m0.004s
sys     0m0.000s
```

Now for Fibonacci we could also easily change the program to an iterative form which would be linear too.

But memoization is a general technique which can be employed in a variety of situations to improve performance. There is even a Perl package to support it.

Performance Improvement Example - Fibonacci

`/home/cs2041/public_html/code/performance/fib2.pl`

```
#!/usr/bin/perl -w
use Memoize;
sub fib($);
memoize('fib');
printf "fib(%d) = %d\n", $_, fib($_) foreach @ARGV;
sub fib($) {
    my ($n) = @_;
    return 1 if $n < 3;
    return fib($n-1) + fib($n-2);
}
```